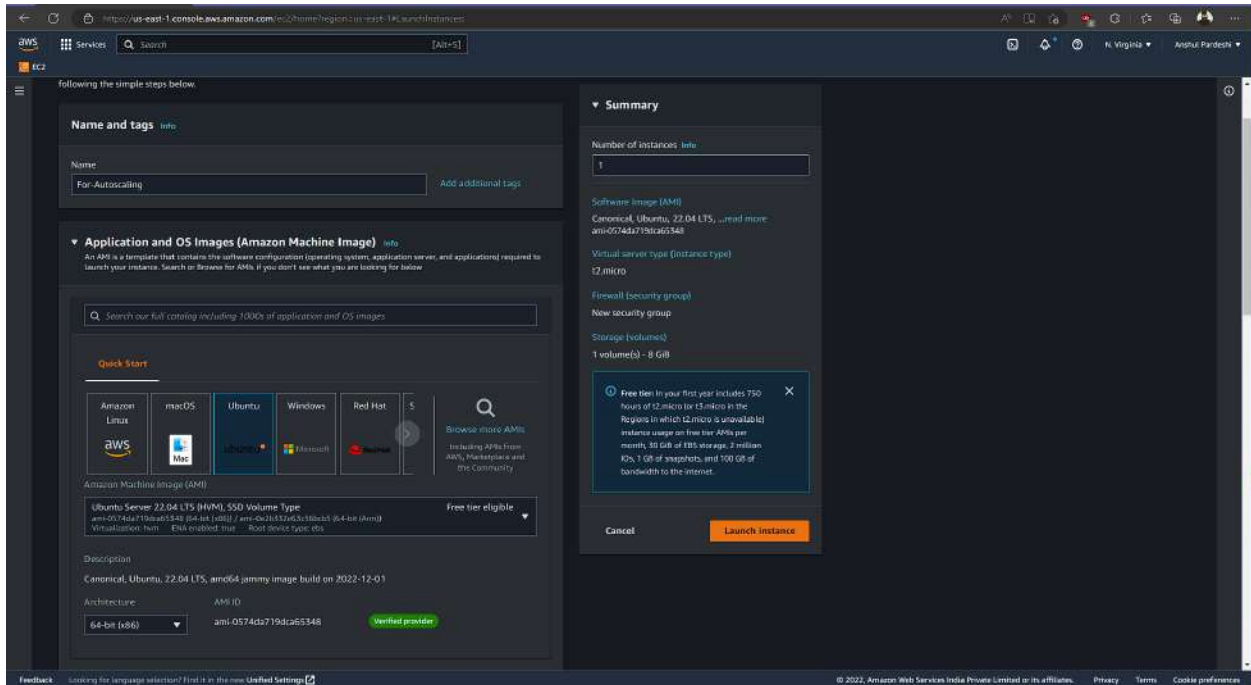


## Module-2: Auto Scaling Assignment - 2


You have been asked to:

1. Create a Web Server AMI with Apache 2 server running in it
2. Create a Launch Configuration with this AMI
3. Use this Launch Configuration to create an Auto Scaling group with 1 minimum and 3 maximum instances.

Create an EC2 instance allowing all traffic of http,ssh and https.



Select the instance created and connect to it.



Connect to that instance and use following commands:

Sudo su

Sudo apt-get update

Sudo apt-get install apache2 -y

cd /var/www/html

rm index.html

nano index ,html and write the Text that you want to be displayed on webserver.



## Give image a name and let other settings to be default

Instance ID: i-0492a96b20263577a (For-AutoScaling)

Image name: For-AutoScaling

Image description - optional: test

No reboot: ☒ Enable

Instance volumes

Volume type	Device	Snapshot	Size	Volume type	IOPS	Throughput	Delete on termination	Encrypted
EBS	/dev/sda1	Create new snapshot from...	8	EBS Standard Provisioned I/O	300	128 MB/s	<input checked="" type="checkbox"/> Enable	<input checked="" type="checkbox"/> Enable

Add volume

During the image creation process, Amazon EC2 creates a snapshot of each of the above volumes.

Tags - optional

A tag is a label that you assign to an AWS resource. Each tag consists of a key and an optional value. You can use tags to search and filter your resources or track your AWS costs.

☒ Tag image and snapshots together  
Tag the image and the snapshots with the same tag.

☐ Tag image and snapshots separately  
Tag the image and the snapshots with different tags.

No tags associated with the resource.

Add new tag

Cancel Create image

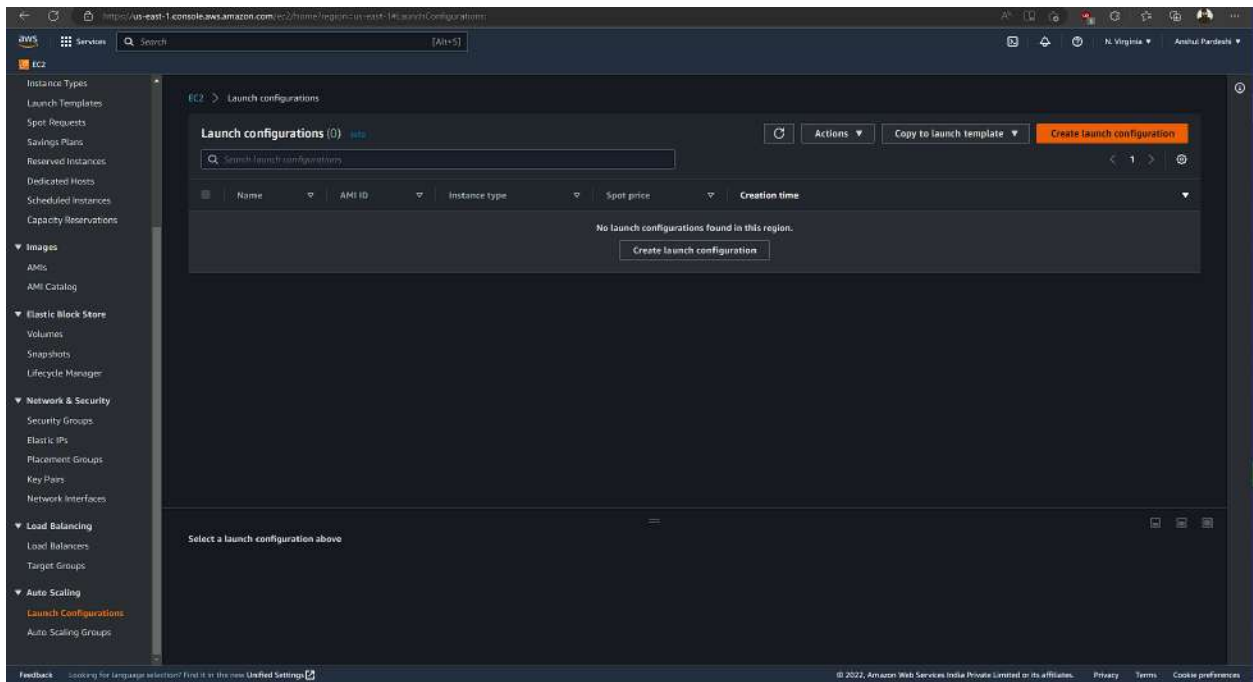
## Check if the image is created.

Amazon Machine Images (AMIs) (1) Info

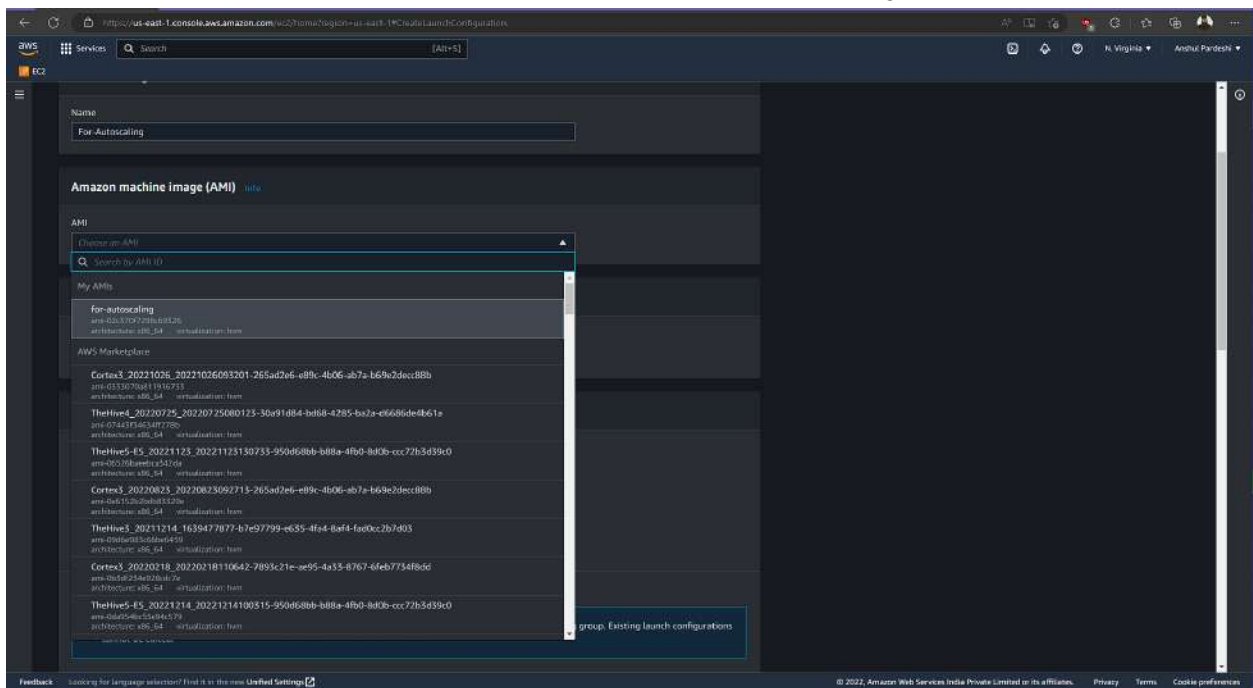
Owned by me Find AMI by attribute or tag

Name	AMI ID	AMI name	Source	Owner	Visibility	Status	Creation date
-	ami-0d5515677e69e0d31	For-AutoScaling	994272653862/For-AutoScaling	994272653862	Private	Pending	2022/12/17 20:09 GMT+5:30

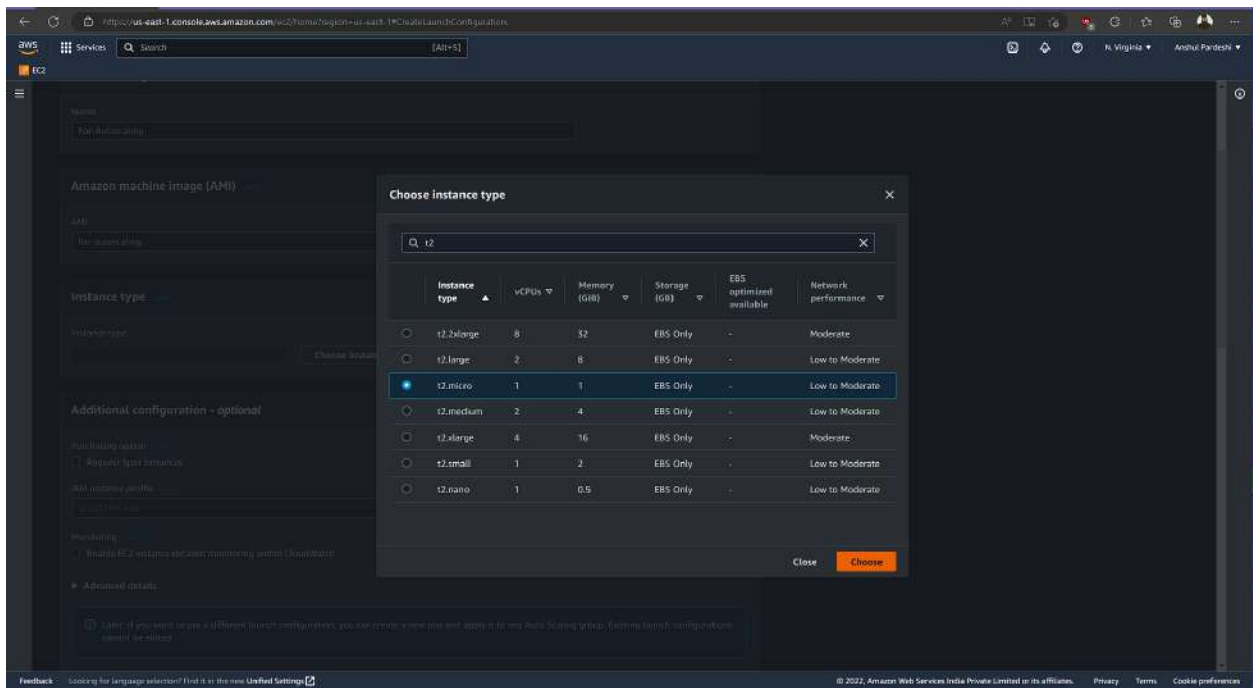
Now let's create Launch Configuration.



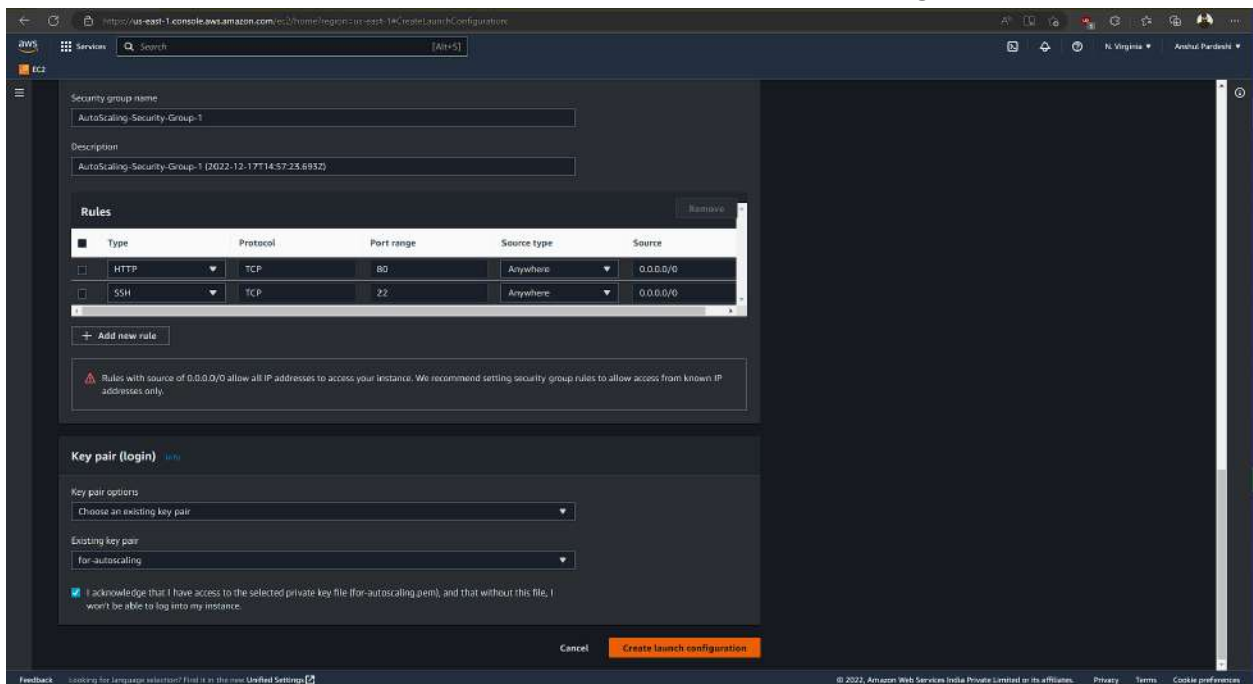
Give a name to it and select the same AMI that we just created.



## Select t2.micro instance for this case



## Allow ssh and http traffic and create Launch Configuration.





## Now let's create Autoscaling group.

The screenshot shows the Amazon EC2 Auto Scaling console. The main heading is "Amazon EC2 Auto Scaling helps maintain the availability of your applications". Below this, a diagram titled "How it works" illustrates an "Auto Scaling group" containing four EC2 instances. Brackets indicate the "Minimum size", "Desired capacity", and "Maximum size" of the group. A "Scale out as needed" label points to the instances. To the right, there are sections for "Create Auto Scaling group", "Pricing", and "Getting started". The "Create Auto Scaling group" section includes a button labeled "Create Auto Scaling group". The "Pricing" section states that Amazon EC2 Auto Scaling features have no additional fees beyond the service fees for Amazon EC2, CloudWatch, and other AWS resources. The "Getting started" section includes links for "What is Amazon EC2 Auto Scaling?", "Getting started with Amazon EC2 Auto Scaling", "Set up a scaled and load-balanced application", and "FAQ".

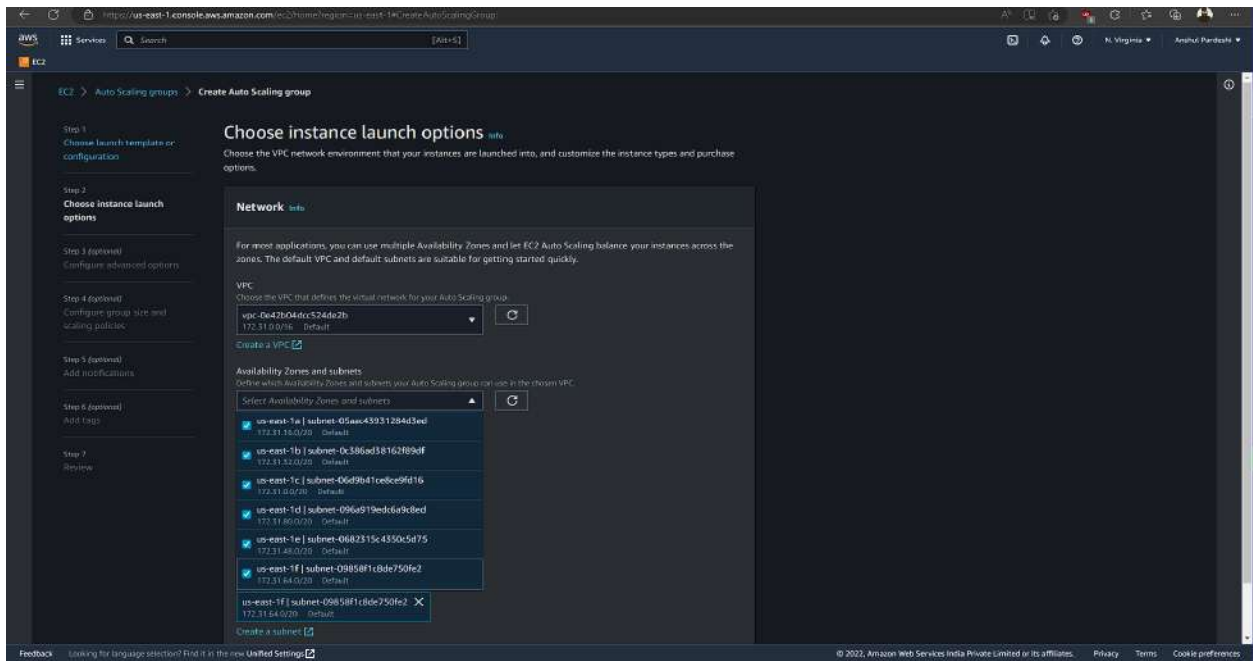
## Give a name to it and select the Launch Configuration that we just created.

The screenshot shows the "Choose launch template or configuration" page in the Amazon EC2 Auto Scaling console. The page is divided into two main sections: "Name" and "Launch configuration". The "Name" section has a text input field for the "Auto Scaling group name" with a placeholder "For-Assignment". The "Launch configuration" section has a dropdown menu for "Launch configuration" with a placeholder "For-Assignment". Below the dropdown, there is a table with details for the selected launch configuration:

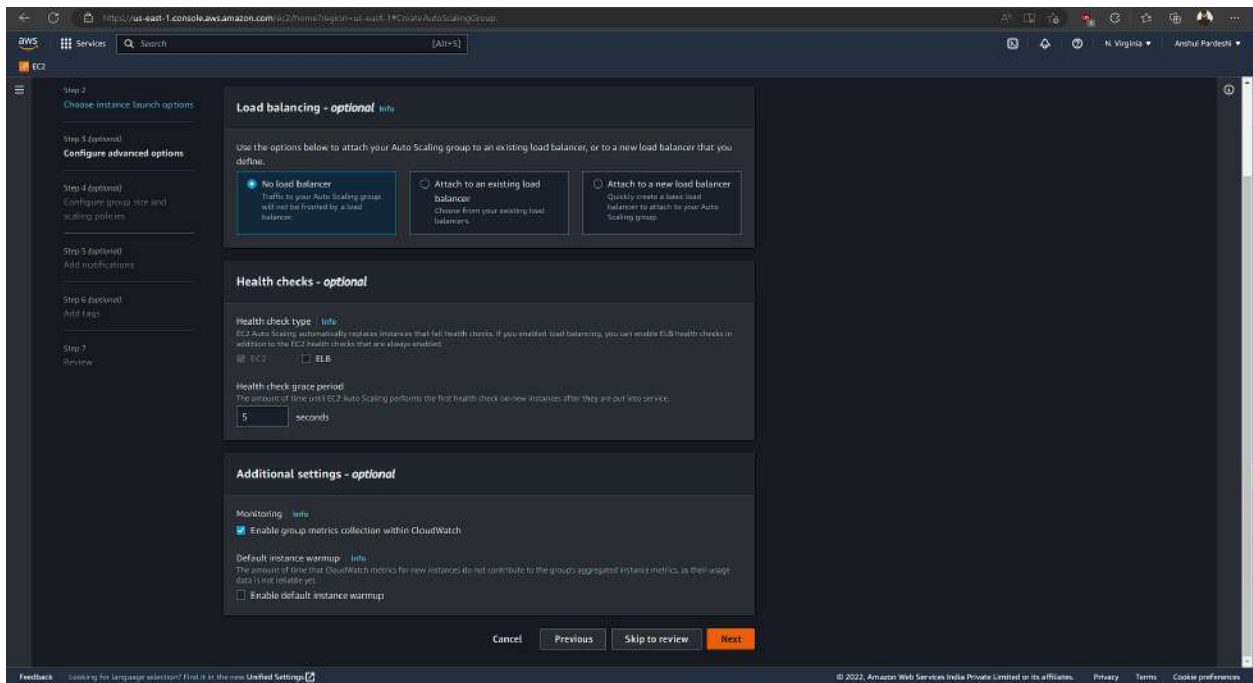
Launch configuration	AMI ID	State created
For-Assignment	ami-02c370f7298c69326	Sat Dec 12 20:55:51 GMT+05:30 (India Standard Time)

Below the table, there is a section for "Security groups" with a link to "Create a launch configuration". At the bottom, there are "Cancel" and "Next" buttons.

Let the VPC be default and then select all AZ's.



Since we are not using load balancer here, lets go with no load balancer.



As mentioned in the assignment, select 1 minimum and 3 maximum instances.

The screenshot shows the 'Create Auto Scaling' wizard in the AWS Management Console, specifically Step 4: 'Configure group size and scaling policies'. The left sidebar lists the steps: Step 2 (Choose instance launch options), Step 3 (Configure advanced options), Step 4 (Configure group size and scaling policies), Step 5 (Add notifications), Step 6 (Add tags), and Step 7 (Review). The main content area is divided into three sections: 'Group size - optional', 'Scaling policies - optional', and 'Instance scale-in protection - optional'. In the 'Group size' section, 'Desired capacity' is set to 1, 'Minimum capacity' is set to 1, and 'Maximum capacity' is set to 3. In the 'Scaling policies' section, the 'None' radio button is selected under 'Choose whether to use a scaling policy to dynamically resize your Auto Scaling group to meet changes in demand.' In the 'Instance scale-in protection' section, the 'Enable instance scale-in protection' checkbox is unchecked. At the bottom, there are buttons for 'Cancel', 'Previous', 'Skip to review', and 'Next'.

Step 2: Choose instance launch options

Step 3 (optional): Configure advanced options

Step 4 (optional): **Configure group size and scaling policies**

Step 5 (optional): Add notifications

Step 6 (optional): Add tags

Step 7: Review

**Group size - optional**

Specify the size of the Auto Scaling group by changing the desired capacity. You can also specify minimum and maximum capacity limits. Your desired capacity must be within the limit range.

Desired capacity: 1

Minimum capacity: 1

Maximum capacity: 3

**Scaling policies - optional**

Choose whether to use a scaling policy to dynamically resize your Auto Scaling group to meet changes in demand.

☐ Target tracking scaling policy

☒ None

**Instance scale-in protection - optional**

Instance scale-in protection

☐ Enable instance scale-in protection

Cancel Previous Skip to review Next

The screenshot shows the 'Create Auto Scaling' wizard in the AWS Management Console, specifically Step 5: 'Add notifications' and Step 6: 'Add tags'. The left sidebar lists the steps: Step 2 (Choose instance launch options), Step 3 (Configure advanced options), Step 4 (Configure group size and scaling policies), Step 5 (Add notifications), Step 6 (Add tags), and Step 7 (Review). The main content area is divided into two sections: 'Step 5: Add notifications' and 'Step 6: Add tags'. In the 'Step 5' section, there is a 'Notifications' section with a 'No notifications' button. In the 'Step 6' section, there is a 'Tags' section with a 'No tags' button. At the bottom, there are buttons for 'Cancel' and 'Create Auto Scaling group'.

Desired capacity: 1 Minimum capacity: 1 Maximum capacity: 3

**Scaling policy**

No scaling policy

**Instance scale-in protection**

Instance scale-in protection

☐ Enable instance scale-in protection

**Step 5: Add notifications**

Notifications

No notifications

**Step 6: Add tags**

Tags (0)

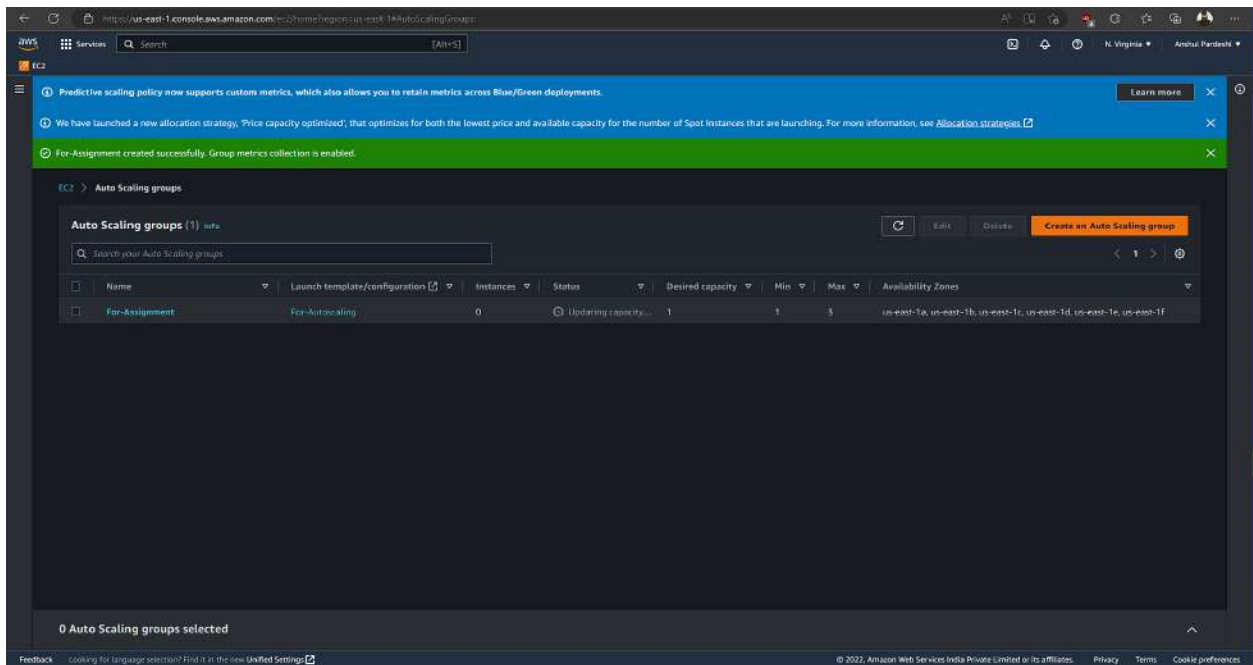
Key Value Tag new instances

No tags

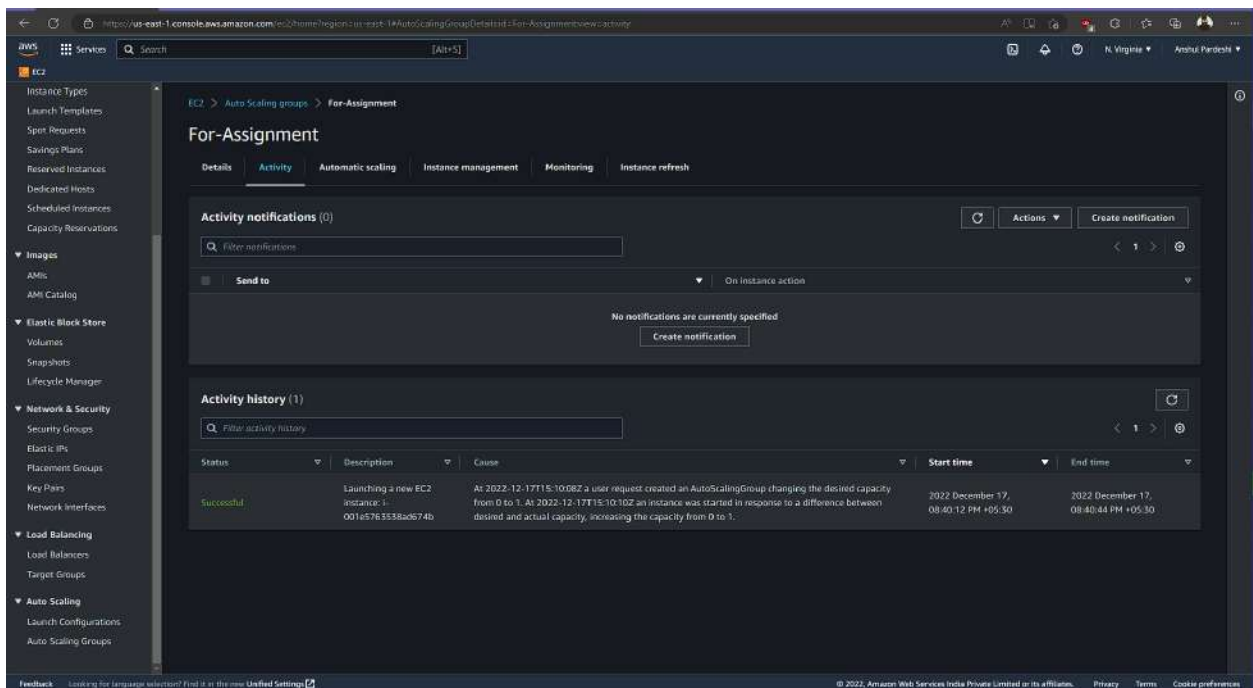
Cancel Create Auto Scaling group



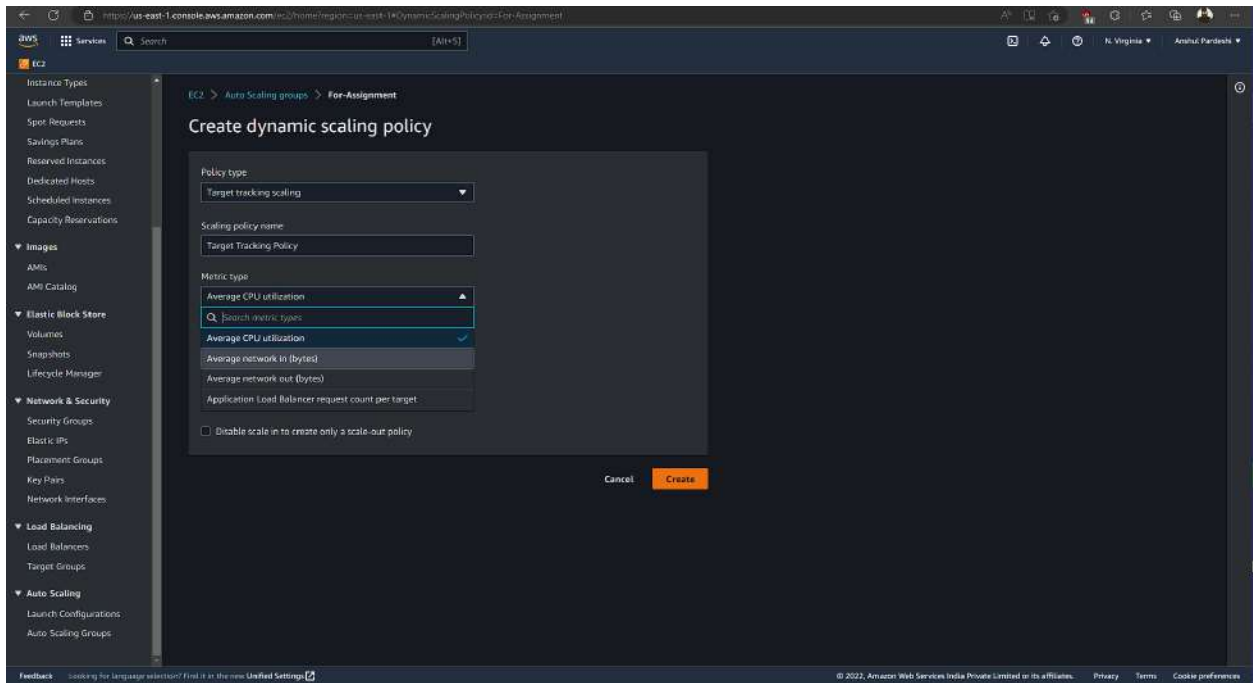
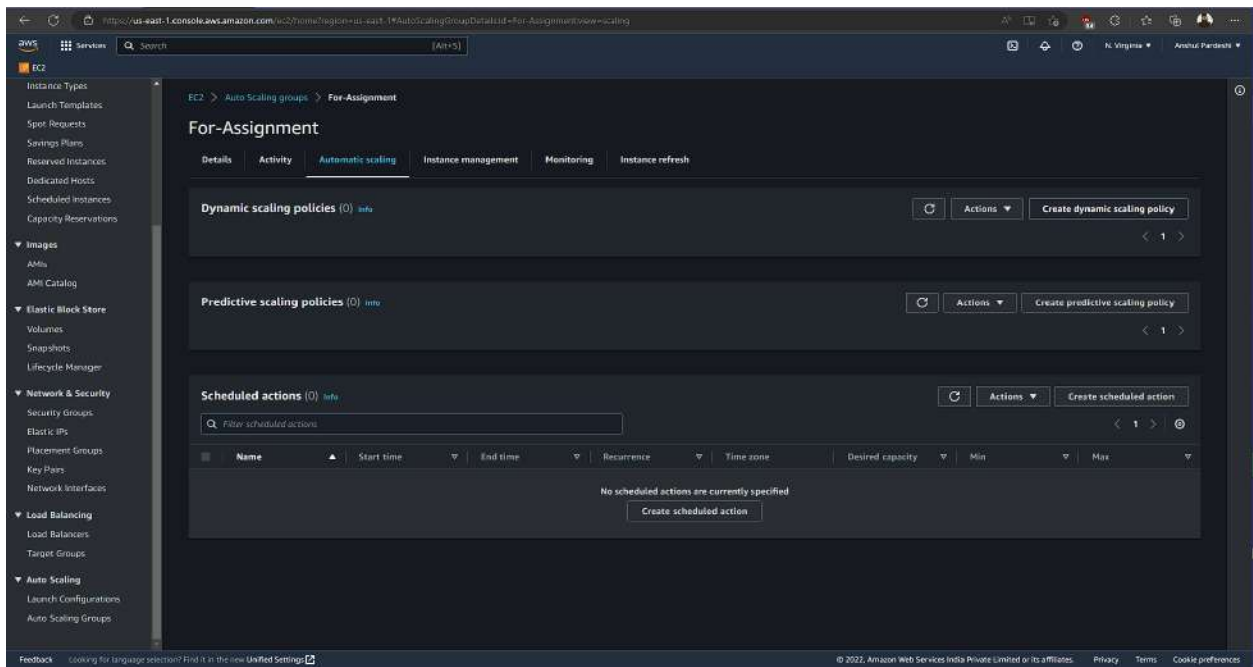
## Keep rest things default and launch Autoscaling Group.



## Open created autoscaling group and on activity you can see number of instances running.



Now to get result faster, let's assign dynamic scaling policy.



Set metric as Avg Network in (bytes) Set Target value as less as possible to get faster results.

The screenshot shows the AWS Management Console interface for creating a dynamic scaling policy. The left sidebar contains navigation links for various AWS services. The main content area is titled 'Create dynamic scaling policy' and includes the following fields:

- Policy type:** Target tracking scaling
- Scaling policy name:** Target Tracking Policy
- Metric type:** Average network in (bytes)
- Target value:** 2
- Instances need:** 300 seconds warm up before including in metric
- ☐ Datable scale in to create only a scale-out policy

Buttons for 'Cancel' and 'Create' are located at the bottom right of the dialog.

The screenshot shows the 'For-Assignment' page for a dynamic scaling policy. The page has tabs for 'Details', 'Activity', 'Automatic scaling', 'Instance management', 'Monitoring', and 'Instance refresh'. A green notification bar at the top states: 'Dynamic scaling policy created or edited successfully.' Below this, the 'Dynamic scaling policies (1/1)' section displays the details of the 'Target Tracking Policy':

- Policy name:** Target Tracking Policy
- Policy type:** Target tracking scaling
- Status:** Enabled
- Target value:** As required to maintain Average Network in at 2
- Warmup:** Add or remove capacity units as required
- Cooldown:** 300 seconds to warm up before including in metric
- Cooldown type:** Enabled

Buttons for 'Actions' and 'Create dynamic scaling policy' are visible. Below this, the 'Predictive scaling policies (0)' section is also shown.

**This is what we get after copying public ip of instance and pasting it in url.  
Now refresh this too many times to overload website to test autoscaling group.**



Now in autoscalinggroup in activity you can see multiple instances being deployed by autoscaling group.

The screenshot shows the AWS Management Console interface for an Auto Scaling Group (ASG) named 'For-Assignment'. The 'Activity' tab is selected, displaying a list of activity notifications and a detailed activity history table.

**Activity notifications (0)**

**Activity history (5)**

Status	Description	Cause	Start time	End time
WaitingForInstanceWarmup	Launching a new EC2 instance i-0516e3507385044f	At 2022-12-17T15:51:15Z a monitor alarm TargetTracking-For-Assignment-AlarmHigh-6d7ed99-cfb-4a5a-96ba-9a51f2602486 in state ALARM triggered policy Target Tracking Policy changing the desired capacity from 1 to 3. At 2022-12-17T15:51:29Z an instance was started in response to a difference between desired and actual capacity, increasing the capacity from 1 to 3.	2022 December 17, 09:21:31 PM +05:30	
WaitingForInstanceWarmup	Launching a new EC2 instance i-0b8b0b01f25f69a22	At 2022-12-17T15:51:15Z a monitor alarm TargetTracking-For-Assignment-AlarmHigh-6d7ed99-cfb-4a5a-96ba-9a51f2602486 in state ALARM triggered policy Target Tracking Policy changing the desired capacity from 1 to 3. At 2022-12-17T15:51:29Z an instance was started in response to a difference between desired and actual capacity, increasing the capacity from 1 to 3.	2022 December 17, 09:21:31 PM +05:30	
Successful	Launching a new EC2 instance i-001e5763538ad674b	At 2022-12-17T15:10:08Z a user request created an AutoScalingGroup changing the desired capacity from 0 to 1. At 2022-12-17T15:10:10Z an instance was started in response to a difference between desired and actual capacity, increasing the capacity from 0 to 1.	2022 December 17, 08:40:12 PM +05:30	2022 December 17, 08:40:44 PM +05:30

The screenshot shows the AWS Management Console interface for the same Auto Scaling Group (ASG) 'For-Assignment'. The 'Instance management' tab is selected, displaying a list of instances and sections for lifecycle hooks and warm pools.

**Instances (3)**

Instance ID	Lifecycle	Instance type	Weighted capacity	Launch template	Availability Zone	Health status	Protected from
i-001e5763538ad674b	InService	t2.micro	-	For-AutoScaling	us-east-1b	Healthy	
i-0516e3507385044f	InService	t2.micro	-	For-AutoScaling	us-east-1d	Healthy	
i-0b8b0b01f25f69a22	InService	t2.micro	-	For-AutoScaling	us-east-1e	Healthy	

**Lifecycle hooks (0)**

**Warm pool**